


Time Indexed Models

Spotting the Problem

Our KDE estimator is learning from all the training data

This means it will learn from both these series:

```
In [5]: plt.figure(figsize=figsize)
plt.plot(wdata.iloc[0], label='first window')
plt.plot(wdata.iloc[1], label='second window')
plt.legend()
plt.tight_layout()
```



Spotting the Problem

Let us consider the first two window applications

- In the first window, the observations are x_0, x_1 and so on
 - In the second window, the observations are x_1, x_2 and so on
- x_0 is number of taxis as 00:00, x_1 at 00:30, and so on
- Hence, the first observation in the first window corresponds to 00:00
 - ...But in the second window corresponds to 00:30

Our estimator learns a distribution for the **observations**:

- Moving the window forward changes "who is who"
- We learn the distribution of x_0 (and its correlations) **multiple times!**

The learning problem is still well defined, but also **very complex**

We will see how to improve things by adding **time information**

Exploiting Time

We have mentioned two ways to exploit time information in KDE:

The first approach consists in making time (of the day) an additional variable:

- Say \mathbf{x} is the vector of our observations and z is the time
- Then we learn an estimator for $P(\mathbf{x}, z)$ and one for $P(z)$
- So that we can compute $P(\mathbf{x} \mid z) = P(\mathbf{x}, z)/P(z)$

In our case $P(z)$ is trivial (uniform)

This is what we did in the previous lecture:

- We managed to use the time information
- ...But we had to add one dimension

This makes the estimation problem more complex

- ...And we are having enough trouble as it is

Exploiting Time

We have discussed two ways to exploit time information in KDE:

The second approach consists in **learning many KDE estimators**:

- Each KDE estimator is **specialized for a given time** (e.g. 00:00, 00:30, 01:00...)
- In other words: our estimator becomes an **ensemble**
- Whole estimation = time-based switch + time-specialized KDE

Let us see the properties of this approach

- The whole model is **more complex** and the approach requires **discrete time**
- Each KDE estimator works with **smaller amounts of data**
- ...But the individual density estimation problems are **easier**!

Let us investigate this approach

Learning a 23:30 Estimator

Let us make a test by learning an estimator for 23:30

First, we separate the training data:

```
In [6]: wdata_tr = wdata[wdata.index < train_end]
wdata_tr.head()
```

Out [6]:

	0	1	2	3	4	5	6	7	8	9	...	38	
timestamp													
2014-07-01 23:30:00	0.357028	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	...	0.883252	0.8
2014-07-02 00:00:00	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	0.082804	...	0.819939	0.7
2014-07-02 00:30:00	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	0.082804	0.143680	...	0.753136	0.6
2014-07-02 01:00:00	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	0.082804	0.143680	0.214862	...	0.671452	0.5
2014-07-02 01:30:00	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	0.082804	0.143680	0.214862	0.363448	...	0.770454	0.8

5 rows × 48 columns

Learning a 23:30 Estimator

Let us make a test by learning an estimator for 23:30

Then, we focus on sequences corresponding to 23:30

```
In [7]: wdata_tr_2330 = wdata_tr.iloc[0::48]
wdata_tr_2330.head()
```

Out [7]:

	0	1	2	3	4	5	6	7	8	9	...	38
timestamp												
2014-07-01 23:30:00	0.357028	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	...	0.883252
2014-07-02 23:30:00	0.440194	0.327429	0.249267	0.194811	0.158694	0.119646	0.098541	0.083462	0.084615	0.081816	...	0.819379
2014-07-03 23:30:00	0.416357	0.347743	0.277088	0.233694	0.191815	0.144306	0.107661	0.097060	0.103579	0.101307	...	0.987225
2014-07-04 23:30:00	0.513318	0.473941	0.412702	0.373391	0.328581	0.276693	0.237053	0.216574	0.186251	0.147302	...	0.493827
2014-07-05 23:30:00	0.578672	0.533006	0.475455	0.412702	0.362361	0.301287	0.263721	0.233629	0.210944	0.145557	...	0.598624

5 rows × 48 columns

Learning a 23:30 Estimator

Then we proceed as usual

We choose a bandwidth:

```
In [8]: grid = GridSearchCV(KernelDensity(kernel='gaussian'),  
                             {'bandwidth': np.linspace(0.01, 0.1, 20)}, cv = 5)  
grid.fit(wdata_tr_2330)  
grid.best_params_
```

```
Out[8]: {'bandwidth': 0.03368421052631579}
```

...And we train the estimator:

```
In [9]: h = grid.best_params_['bandwidth']  
kde_2330 = KernelDensity(kernel='gaussian', bandwidth=h)  
kde_2330.fit(wdata_tr_2330);
```

- We just need to repeat this for each (meaningful) time value

Learning the Ensemble

In our case, all unique hour values are meaningful

```
In [10]: day_hours = wdata_tr.index.hour + wdata_tr.index.minute / 60
day_hours = day_hours.unique()
print(day_hours)

Float64Index([23.5,  0.0,  0.5,  1.0,  1.5,  2.0,  2.5,  3.0,  3.5,  4.0,  4.5,
              5.0,  5.5,  6.0,  6.5,  7.0,  7.5,  8.0,  8.5,  9.0,  9.5, 10.0,
              10.5, 11.0, 11.5, 12.0, 12.5, 13.0, 13.5, 14.0, 14.5, 15.0, 15.5,
              16.0, 16.5, 17.0, 17.5, 18.0, 18.5, 19.0, 19.5, 20.0, 20.5, 21.0,
              21.5, 22.0, 22.5, 23.0],
              dtype='float64', name='timestamp')
```

- `unique` in pandas returns a `series` with all unique values
- It does not matter how we measure time
- ...We only care about having 48 discrete steps

Learning the Ensemble

Then we learn 48 specialized estimators

```
In [11]: kde = {}  
         for hidx, hour in enumerate(day_hours):  
             tmp_data = wdata_tr.iloc[hidx::48]  
             kde[hour] = KernelDensity(kernel='gaussian', bandwidth=h)  
             kde[hour].fit(tmp_data)
```

- For each unique time value, we separate a subset of the training data
- Then we build and learn a KDE estimator

We chose to store everything in a dictionary:

```
In [12]: print(str(kde)[:256], '...}')  
  
{23.5: KernelDensity(bandwidth=0.03368421052631579), 0.0: KernelDensity(bandwidth=0.03368421052631579), 0.5: KernelDensity(bandwidth=0.03368421052631579), 1.0: KernelDensity(bandwidth=0.03368421052631579), 1.5: KernelDensity(bandwidth=0.03368421052631579), ...}
```

Generating the Signal

The we can generate the alarm signal

- In a practical implementation we should do this step by step
- ...But for an evaluation purpose it is easier to do it all at once

```
In [13]: ldens_list = []
         for hidx, hour in enumerate(day_hours):
             tmp_data = wdata.iloc[hidx::48]
             tmp_ldens = kde[hour].score_samples(tmp_data)
             tmp_ldens = pd.Series(index=tmp_data.index, data=tmp_ldens)
             ldens_list.append(tmp_ldens)
```

- For each unique time value, we separate a subset of the whole data
- Then we obtain the estimated (log) probabilities

The process is faster than before!

- ...Because each KDE estimator is trained a smaller dataset

Generating the Signal

All signals are stored in a list

- We need to concatenate them all in single `DataFrame`
- Then we can sort all rows by timestamp (it's the index)

```
In [14]: ldens = pd.concat(ldens_list, axis=0)
         ldens = ldens.sort_index()
         signal = -ldens
         signal.head()
```

```
Out[14]: timestamp
         2014-07-01 23:30:00    -113.900870
         2014-07-02 00:00:00    -113.909604
         2014-07-02 00:30:00    -113.909604
         2014-07-02 01:00:00    -113.909604
         2014-07-02 01:30:00    -113.909604
         dtype: float64
```


A suggestion: always do concatenations in a single step in `pandas`

It's way faster than appending `DataFrame` objects one by one

Generating the Signal

Now we can plot out signal:

```
In [15]: nab.plot_series(signal, labels=labels, windows=windows, figsize=figsize)
```




It's way more stable than before!

Effect of the Threshold

Now, as usual, we examine the response surface

```
In [16]: thr_range = np.linspace(20, 200, 100)
cost_range = [cmodel.cost(signal, labels, windows, thr)
              for thr in thr_range]
cost_range = pd.Series(index=thr_range, data=cost_range)
nab.plot_series(cost_range, figsize=figsize)
```



Threshold Optimization

Then we optimize the threshold:

```
In [17]: signal_opt = signal[signal.index < val_end]
labels_opt = labels[labels < val_end]
windows_opt = windows[windows['end'] < val_end]

best_thr, best_cost = nab.opt_thr(signal_opt, labels_opt,
                                windows_opt, cmodel, thr_range)
print(f'Best threshold: {best_thr}, corresponding cost: {best_cost}')
```

```
Best threshold: 43.63636363636364, corresponding cost: 10
```

Let us see the cost on the whole dataset:

```
In [18]: ctst = cmodel.cost(signal, labels, windows, best_thr)
print(f'Cost on the whole dataset {ctst}')
```

```
Cost on the whole dataset 14
```

Considerations

Some considerations and take-home messages

Be mindful of trade-offs

- We wanted to exploit some new information, and we had to choose between:
 - Using all data, but complicating the learning problem
 - Loosing some data, but simplifying the learning problem
- We chose the latter, but perhaps even the other approach had a chance!

Careful with time-dependent models

- Using raw, absolute, time as a feature may lead to **overfitting**
- ...Since every sample has its own time
- So, exploiting time relies on the existend ot meaningful **periods**